

Efficient Knowledge Distillation for Image Classification: A Comprehensive Study

Regina Wan
University of Auckland
New Zealand, Auckland

Abstract—This report investigates the application of knowledge distillation techniques to create efficient image classification models. We explore the process of transferring knowledge from a large, complex teacher model to a smaller, more efficient student model while maintaining competitive accuracy. Our methodology includes comprehensive data preprocessing, model architecture design, and distillation process optimization. Experimental results demonstrate that the distilled student model achieves 92.3% of the teacher model's accuracy while requiring only 30% of the computational resources. The report concludes with a discussion of limitations and potential future research directions in knowledge distillation.

Keywords—knowledge distillation, model compression, image classification, deep learning, efficiency optimization

I. INTRODUCTION

This template, modified in MS Word 2007 and saved as a “Word 97-2003 Document” for the PC, provides authors with most of the formatting specifications needed for preparing electronic versions of their papers. All standard paper components have been specified for three reasons: (1) ease of use when formatting individual papers, (2) automatic compliance to electronic requirements that facilitate the concurrent or later production of electronic products, and (3) conformity of style throughout a conference proceedings. Margins, column widths, line spacing, and type styles are built-in; examples of the type styles are provided throughout this document and are identified in italic type, within parentheses, following the example. Some components, such as multi-leveled equations, graphics, and tables are not prescribed, although the various table text styles are provided. The formatter will need to create these components, incorporating the applicable criteria that follow.

II. LITERATURE REVIEW

Knowledge distillation was first introduced by Hinton et al. [1], who demonstrated that a smaller student network could learn to mimic the behavior of a larger teacher network by training on both ground truth labels and the teacher's softened output probabilities. This foundational work has inspired numerous extensions and improvements:

1. **Response-based distillation:** The original approach [1] focuses on matching the teacher's output distribution using a temperature-scaled softmax function. Subsequent work by Ba and Caruana [2] showed that even shallow networks could learn complex functions when trained to mimic deeper networks.

2. **Feature-based distillation:** Romero et al. [3] introduced FitNets, which extend distillation to intermediate layers by matching feature representations between teacher and student. This approach has proven particularly effective for very deep networks.

3. **Attention transfer:** Zagoruyko and Komodakis [4] demonstrated that transferring attention maps could further improve student performance, as these maps capture important spatial relationships learned by the teacher.

4. **Relational knowledge distillation:** Park et al. [5] proposed distilling relationships between different data samples rather than individual outputs, capturing higher-level structural knowledge.

5. **Dynamic distillation:** You et al. [6] introduced adaptive methods that automatically determine the importance of different knowledge components during training.

6. **Quantitative analyses:** Recent work by Cho and Hariharan [7] has provided empirical evidence about the effectiveness of different distillation techniques across various architectures and datasets.

Despite these advances, several gaps remain in the literature. Most studies focus on standard benchmark datasets, leaving open questions about distillation effectiveness in real-world scenarios with noisy or limited data. Additionally, the interaction between distillation and other model compression techniques (e.g., quantization, pruning) requires further investigation.

III. METHODOLOGY

A. Data Preprocessing

Our experiments used the CIFAR-100 dataset [8], which consists of 60,000 32×32 color images across 100 classes. We applied the following preprocessing steps:

1. **Normalization:** Pixel values were scaled to [0,1] range

2. **Data augmentation:** Random horizontal flips and 4-pixel padding with random crops

3. **Train/validation/test split:** 45,000/5,000/10,000 images respectively

B. Model Architecture

We implemented two model architectures:

1. Teacher Model: A ResNet-50 [9] pretrained on ImageNet, fine-tuned on CIFAR-100

2. Student Model: A custom lightweight CNN with:

- 4 convolutional blocks (Conv2D-BatchNorm-ReLU-MaxPool)
- Final dense layer with softmax activation
- Approximately $10\times$ fewer parameters than the teacher

C. Knowledge Distillation Process

The distillation process combined three loss components:

- Classification loss (L_{cls}): Standard cross-entropy with true labels
- Distillation loss (L_{dist}): KL-divergence between teacher and student outputs
- Feature loss (L_{feat}): Mean squared error between intermediate feature maps

The total loss was computed as:

$$L_{total} = \alpha L_{cls} + \beta L_{dist} + \gamma L_{feat}$$

where α , β , and γ were hyperparameters balancing the loss components.

D. Training Protocol

We trained the student model using:

- Optimizer: Adam with learning rate 0.001
- Batch size: 128
- Temperature parameter (τ): 4 for softening probabilities
- Training epochs: 200 with early stopping
- Loss weights: $\alpha=0.5$, $\beta=0.3$, $\gamma=0.2$ (determined via validation)

IV. RESULTS

A. Model Performance

Table I compares the performance of our models:

TABLE I. MODEL PERFORMANCE COMPARISON

Metric	Teacher Model	Student Model	Standalone Student
Top-1 Accuracy	76.8%	70.9%	65.2%
Top-5 Accuracy	93.4%	89.7%	85.1%
Parameters	23.5M	2.3M	2.3M
Inference Time	42ms	12ms	12ms

The distilled student achieved 92.3% of the teacher's accuracy while using only 9.8% of the parameters and requiring 28.6% of the inference time.

B. Training Dynamics

Fig. 1 shows the training curves for the student model with and without distillation:

Fig. 1. Training accuracy curves comparing distilled vs. standalone student models. (figure caption)

The distilled student converged faster and to a higher final accuracy than the standalone version, demonstrating the effectiveness of knowledge transfer.

C. Computational Efficiency

We measured computational requirements on an NVIDIA V100 GPU:

- Teacher model: 15.7 GFLOPs per inference
- Student model: 1.2 GFLOPs per inference (7.6% of teacher)

The memory footprint reduction was similarly significant, making the student model suitable for deployment on mobile devices.

V. DISCUSSION

Our results confirm that knowledge distillation can effectively compress deep neural networks while preserving most of their predictive power. The distilled student model achieved significantly better performance than its standalone counterpart, demonstrating that the teacher's knowledge was successfully transferred.

Several factors contributed to this success:

- The temperature parameter effectively softened the teacher's output distribution, revealing class relationships beyond hard labels
- Intermediate feature matching helped the student learn useful representations
- The balanced loss function prevented overemphasis on any single knowledge component

However, we identified several limitations:

- The performance gap between teacher and student remains non-negligible
- The distillation process requires significant computational resources for training
- Optimal hyperparameters (α , β , γ , τ) are dataset-dependent and require tuning

Future work could explore:

1. Combining distillation with other compression techniques
2. Developing adaptive methods for automatic hyperparameter tuning
3. Investigating distillation for more complex tasks like object detection

This project significantly enhanced our understanding of model compression techniques and the practical challenges involved in deploying efficient deep learning models.

VI. CONCLUSION

We presented a comprehensive study of knowledge distillation for image classification, demonstrating that careful implementation can yield compact models with minimal accuracy loss. Our distilled student model achieved over 90% of the teacher's accuracy while being dramatically more efficient. These findings have important implications for deploying deep learning models in resource-constrained environments.

Future research directions include investigating distillation for other computer vision tasks, developing more sophisticated knowledge transfer methods, and creating theoretical frameworks to explain distillation's effectiveness. The growing need for efficient AI makes knowledge distillation an increasingly important area of study with substantial practical impact.

REFERENCES

- [1] G. Hinton, O. Vinyals, and J. Dean, "Distilling the knowledge in a neural network," arXiv preprint arXiv:1503.02531, 2015.
- [2] J. Ba and R. Caruana, "Do deep nets really need to be deep?" in *Advances in Neural Information Processing Systems*, 2014, pp. 2654-2662.
- [3] A. Romero et al., "FitNets: Hints for thin deep nets," arXiv preprint arXiv:1412.6550, 2014.
- [4] S. Zagoruyko and N. Komodakis, "Paying more attention to attention: Improving the performance of convolutional neural networks via attention transfer," arXiv preprint arXiv:1612.03928, 2016.
- [5] W. Park et al., "Relational knowledge distillation," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2019, pp. 3967-3976.
- [6] S. You et al., "Drawing early-bird tickets: Towards more efficient training of deep networks," arXiv preprint arXiv:1909.11957, 2019.
- [7] J. H. Cho and B. Hariharan, "On the efficacy of knowledge distillation," in *Proceedings of the IEEE/CVF International Conference on Computer Vision*, 2019, pp. 4794-4802.
- [8] A. Krizhevsky, "Learning multiple layers of features from tiny images," University of Toronto, Tech. Rep., 2009.
- [9] K. He et al., "Deep residual learning for image recognition," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2016, pp. 770-778.